出门问问

Accelerating Deep Learning Inference

Approaches Overview: Concepts and Examples

Shen Li 2018.07.26

Deep Learning Dev Process

出门问问

Task Modeling -
- Input and Output Representation
- Network Architecture Design

Model Training -
- Knowledge Distillation
- Transfer/Multi-task Learning

Model Conversion -
- Model Pruning and Quantization
- Computation Graph Optimization

Model Deployment -
- Optimized Kernel Implementation
- Hardware Acceleration

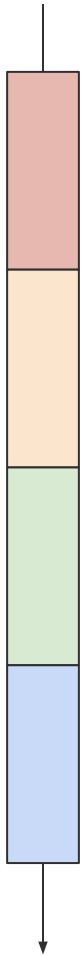

Task Modeling

Efficient Representation for Inputs and Outputs

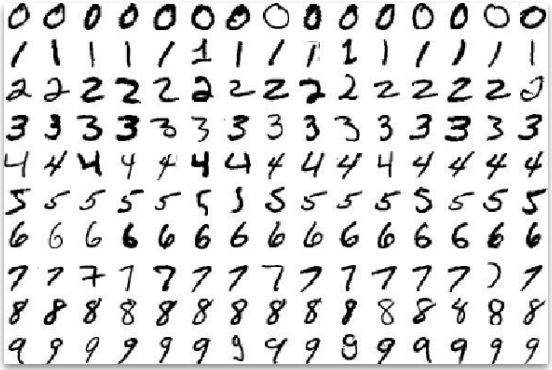

MNIST Handwritten Digits Dataset

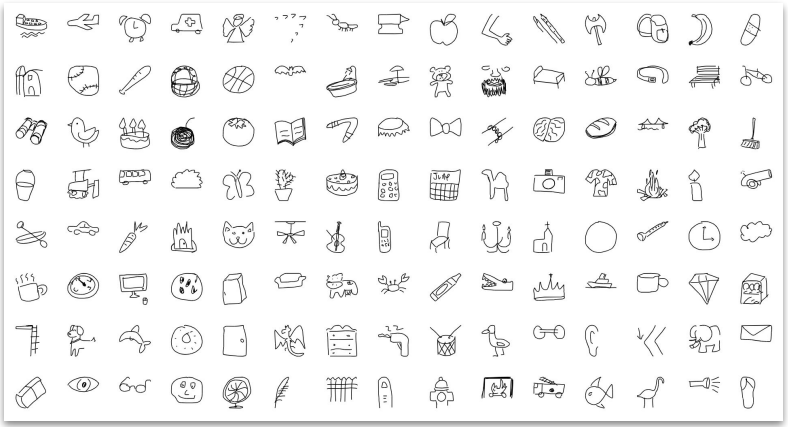

Quick,Draw! Drawing Dataset

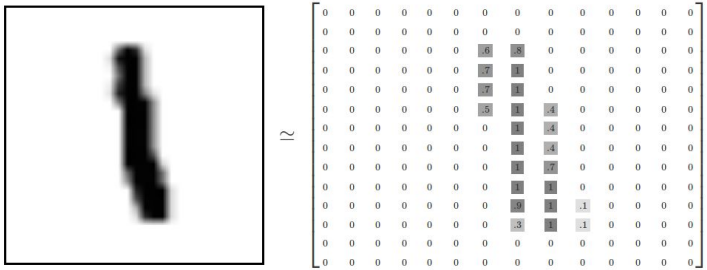

$$\begin{bmatrix} 0 & 0.2 & -0.3 & \dots & -0.7 & \dots \\ 0 & 0.5 & 0.4 & \dots & 0.9 & \dots \\ 1 & 0 & 0 & \dots & 1 & \dots \end{bmatrix}$$

Task Modeling

Efficient Representation for Inputs and Outputs

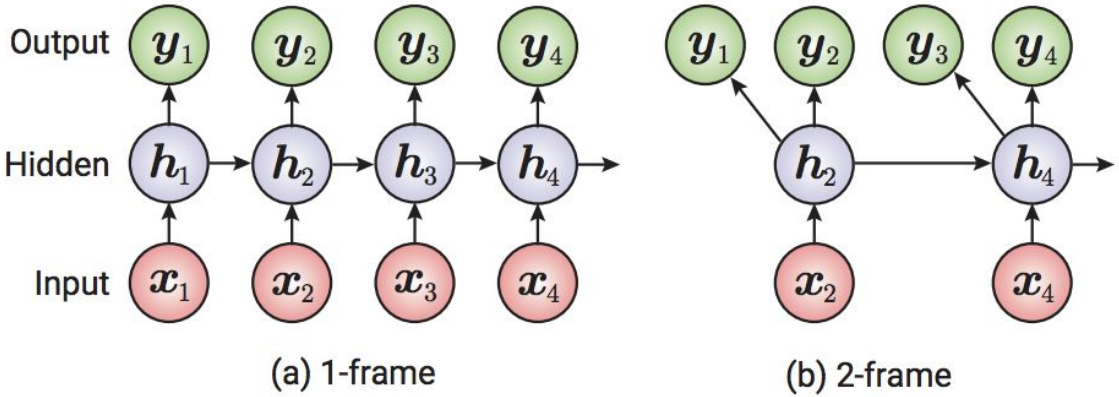

Multi-frame bundled inference for RNN

Zen, Heiga, et al. "Fast, compact, and high quality LSTM-RNN based statistical parametric speech synthesizers for mobile devices." *arXiv preprint arXiv:1606.06061* (2016).

Neural Network Design

Factorized Neural Network

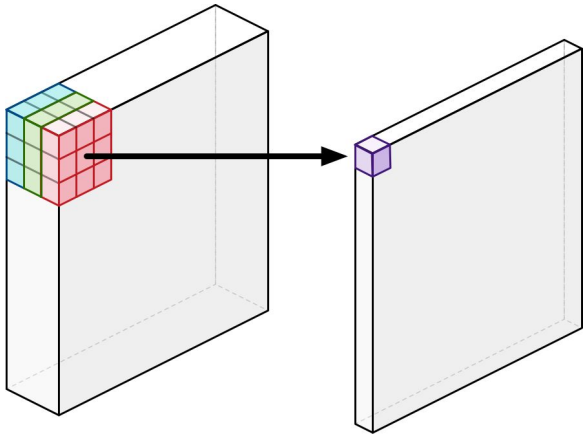

Regular Convolution

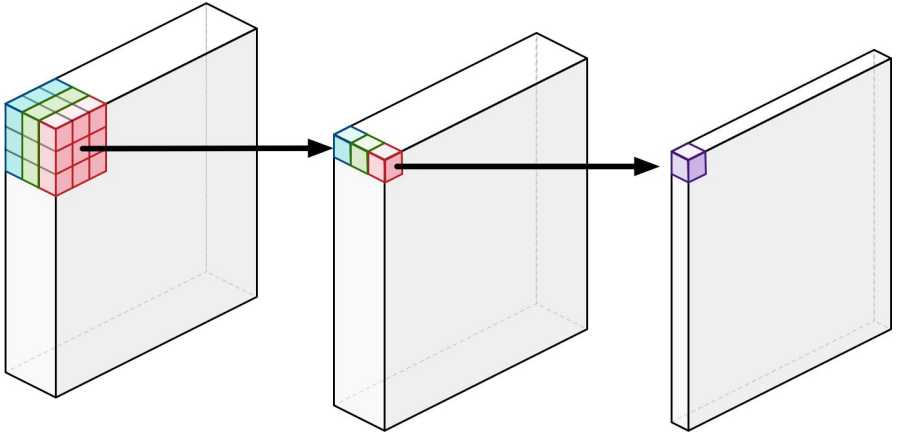

Depthwise Separable Convolution

http://machinethink.net/blog/googles-mobile-net-architecture-on-iphone/

Neural Network Design

Neural Architecture Search

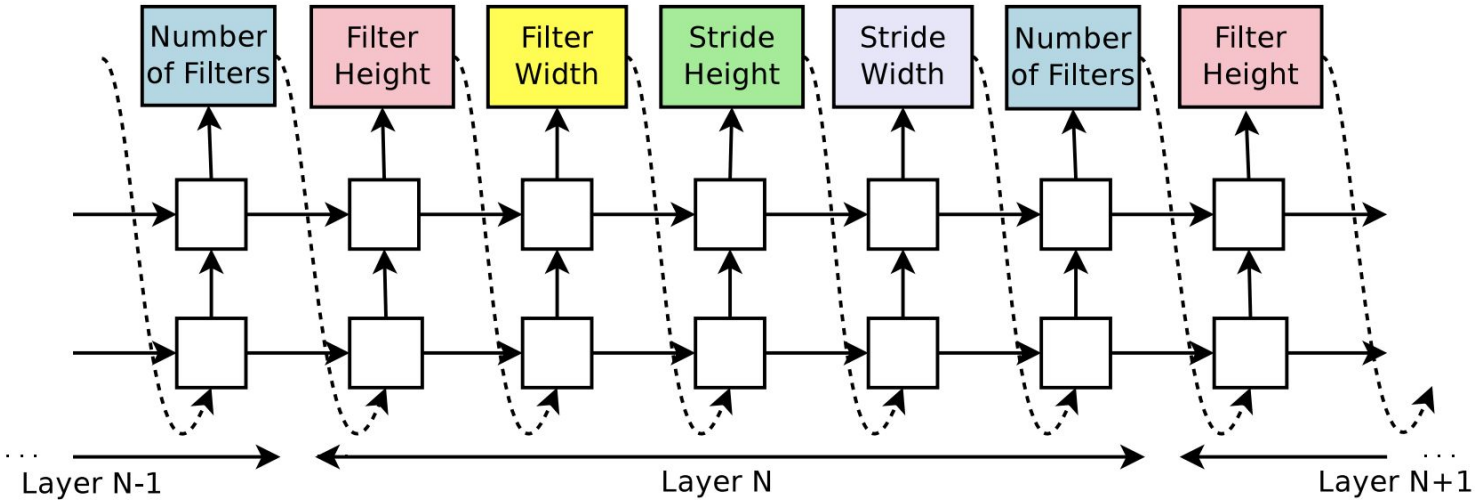

Zoph, Barret, and Quoc V. Le. "Neural architecture search with reinforcement learning." arXiv preprint arXiv:1611.01578 (2016).

Knowledge Distillation

Guided Training Student Model from Teacher Model

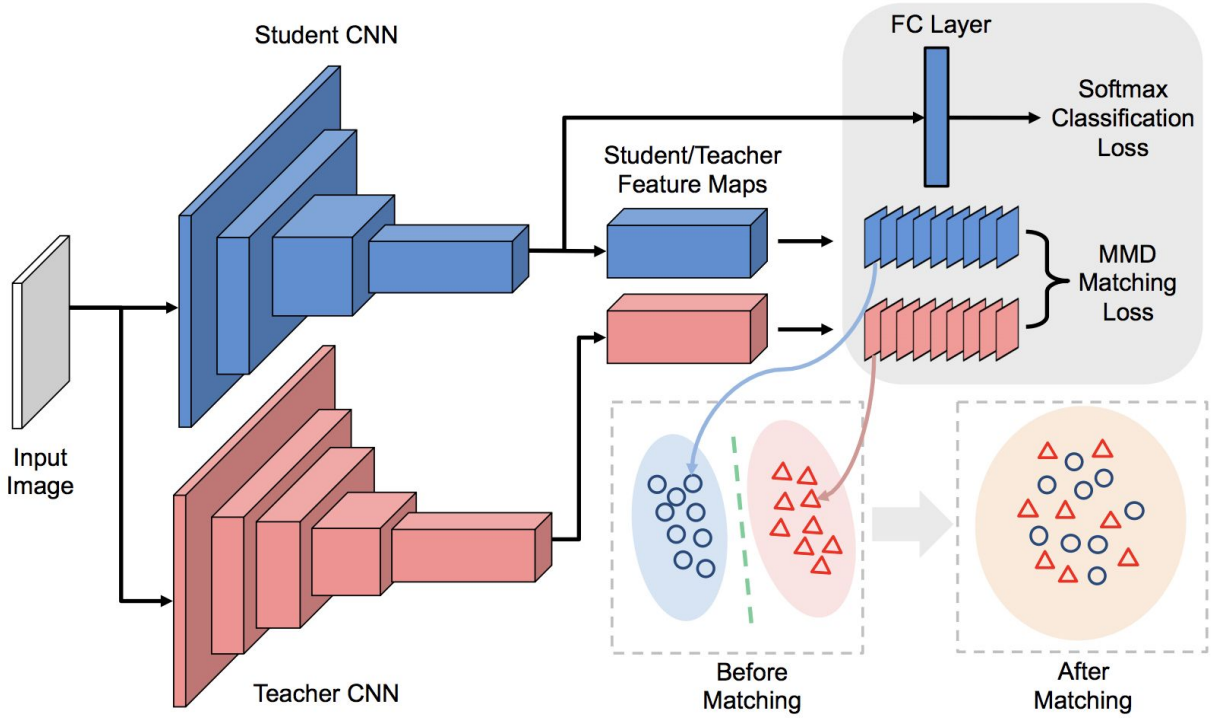

Huang, Zehao, and Naiyan Wang. "Like what you like: Knowledge distill via neuron selectivity transfer." arXiv preprint arXiv:1707.01219 (2017).

Model Compression

Neural Network Model Pruning

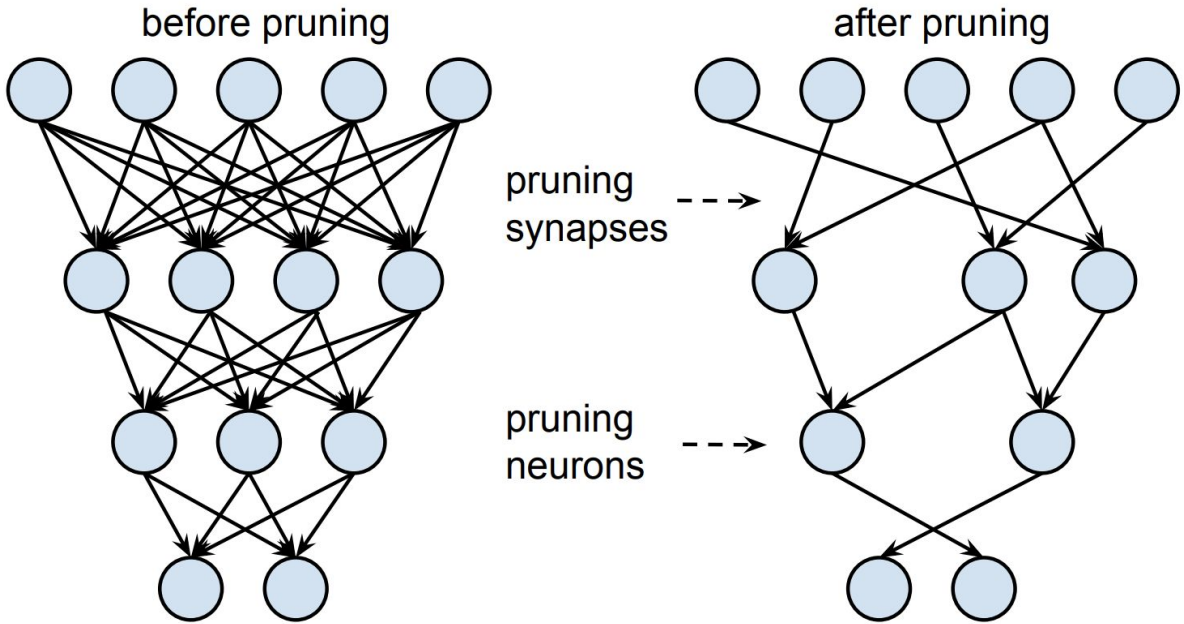

Han, Song, et al. "Learning both weights and connections for efficient neural network." Advances in neural information processing systems. 2015.

Model Compression

Model Quantization

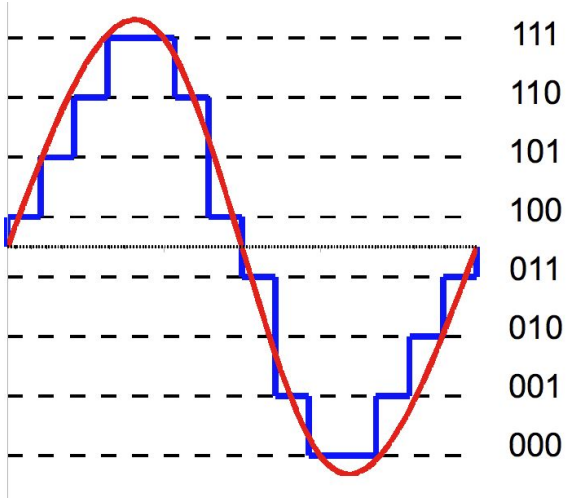

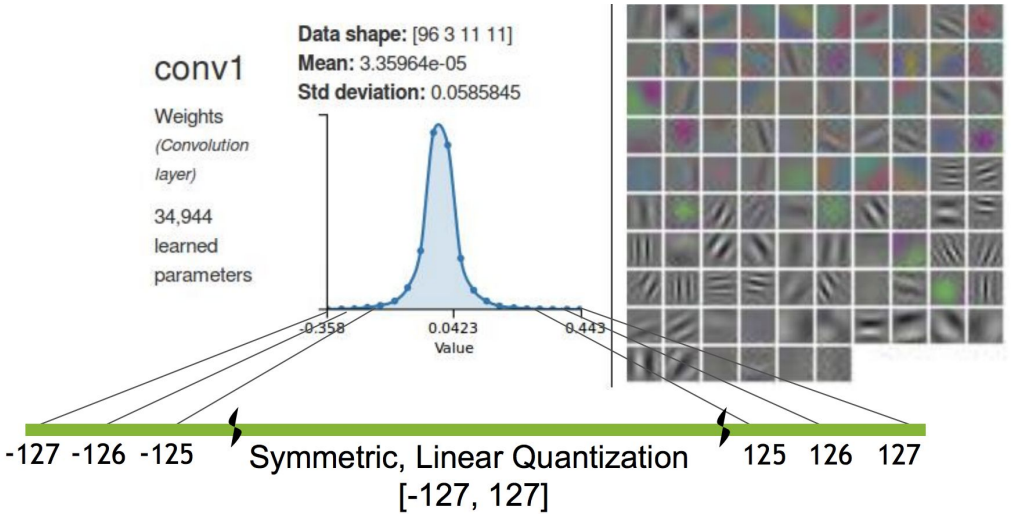

Model Compression

Model Quantization

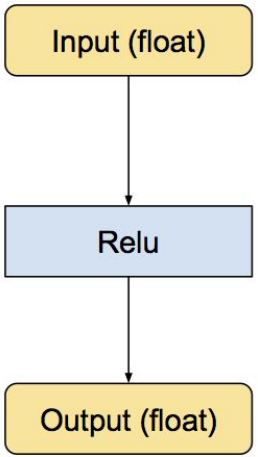

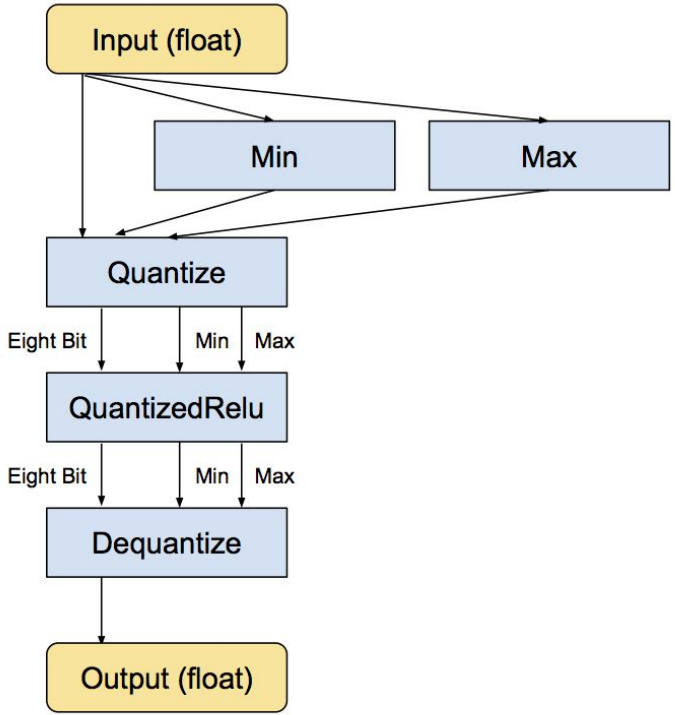

https://petewarden.com/2016/05/03/how-to-quantize-neural-networks-with-tensorflow/

Model Compression

Model Quantization

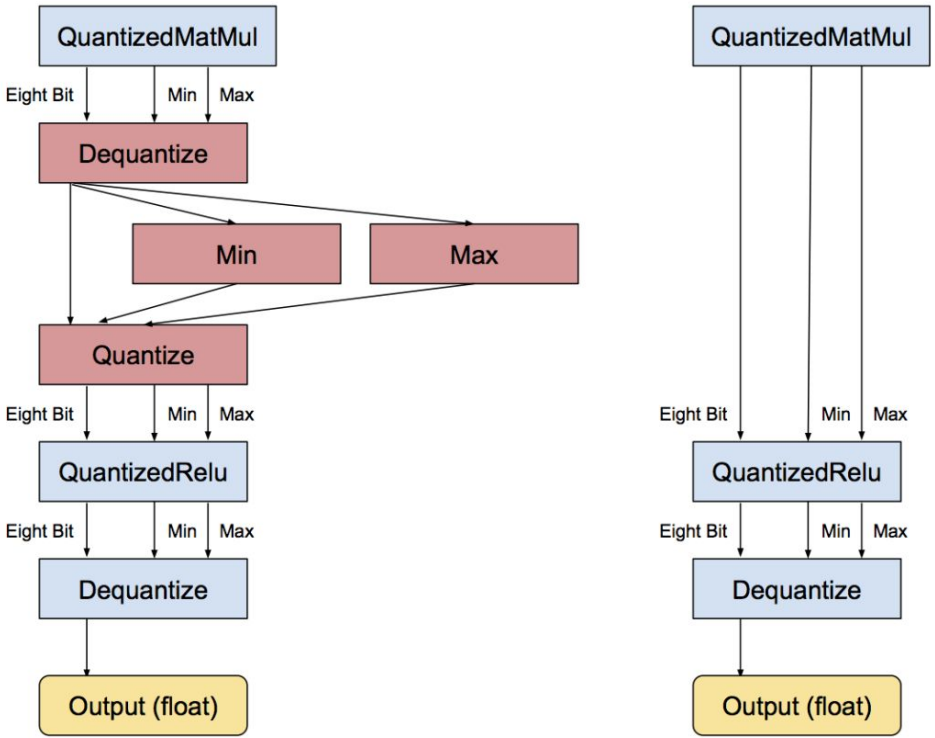

https://petewarden.com/2016/05/03/how-to-quantize-neural-networks-with-tensorflow/

Computation Graph Optimization

Operation Fusion

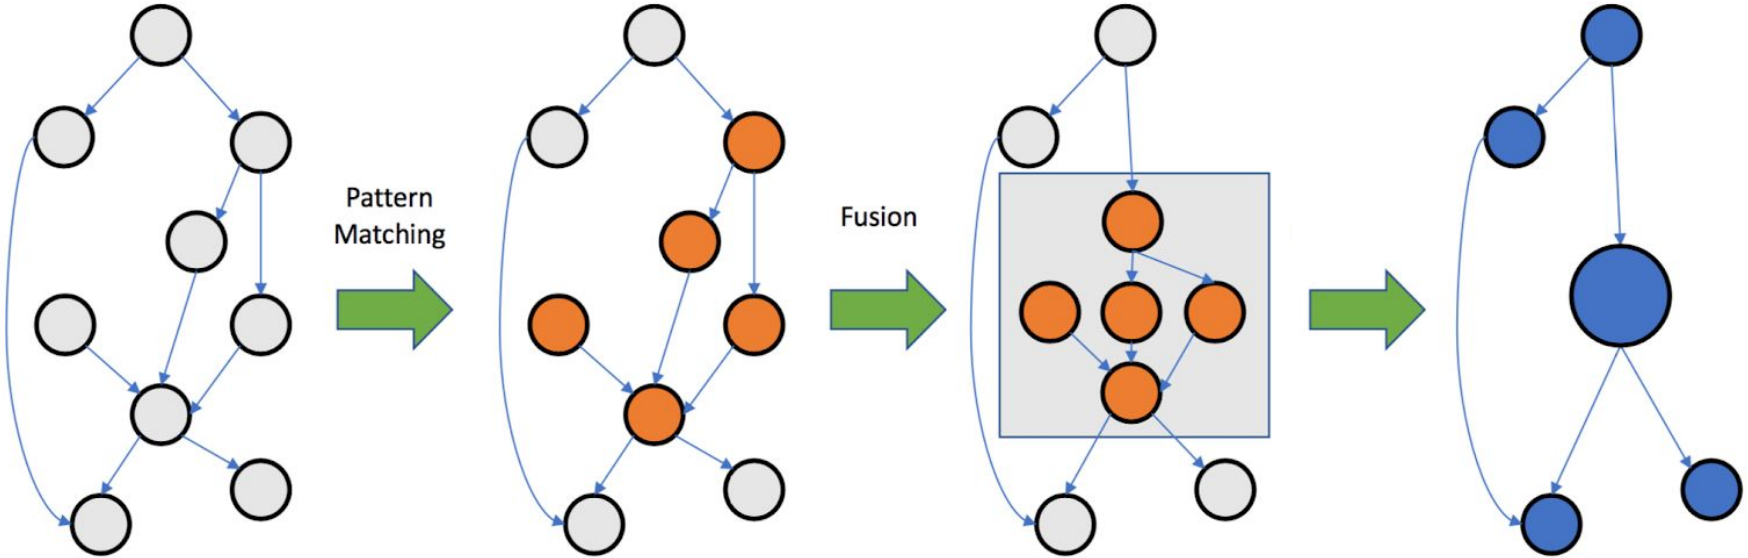

Computation Graph Optimization

Operation Fusion

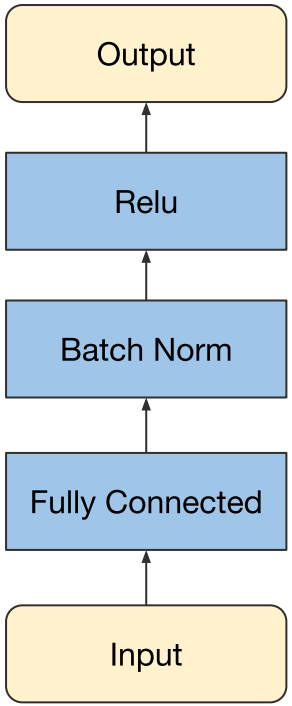

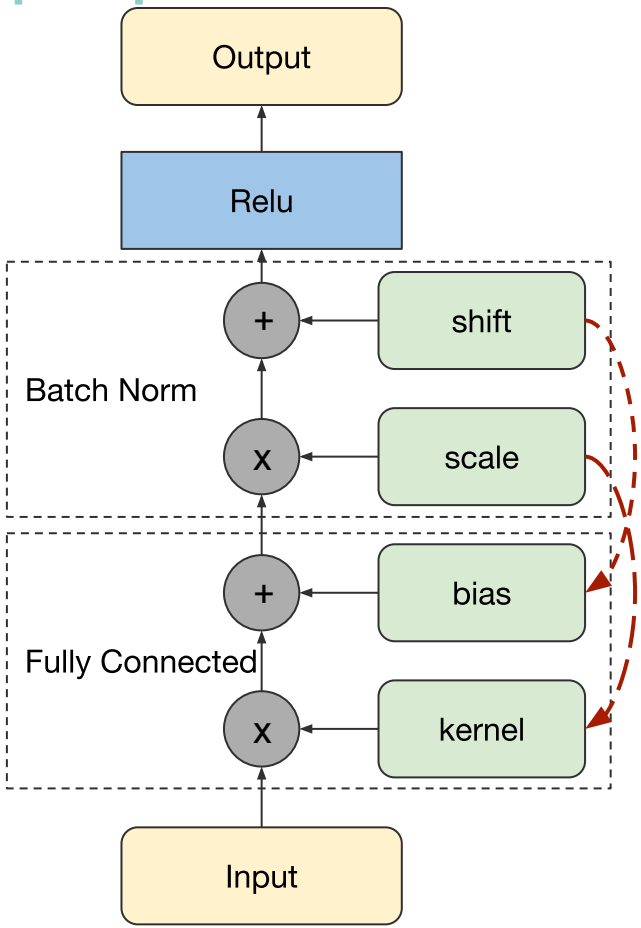

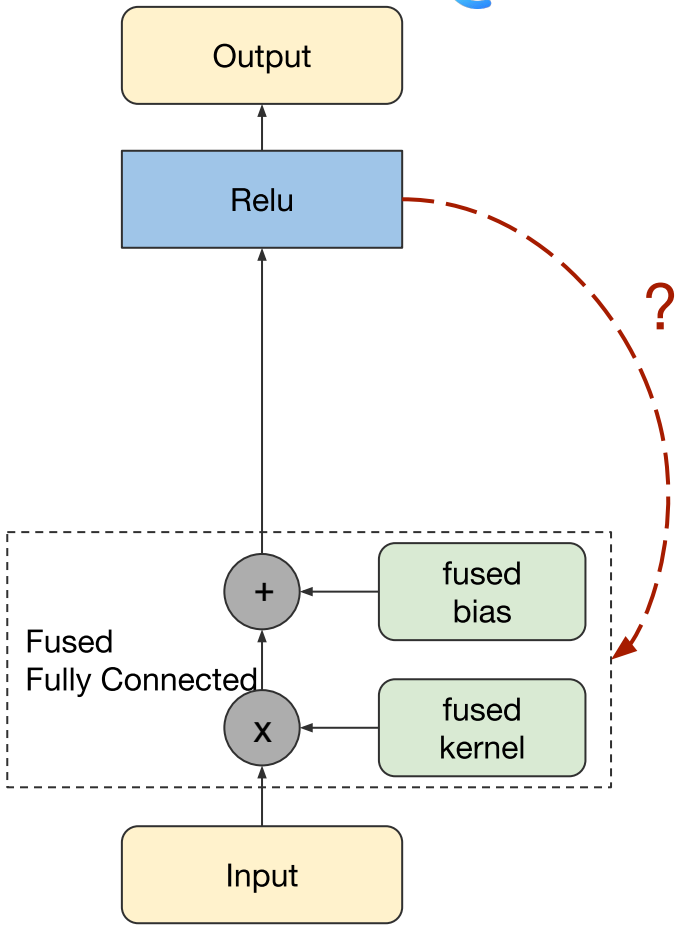

Computation Graph Optimization

AOT & JIT Compilation

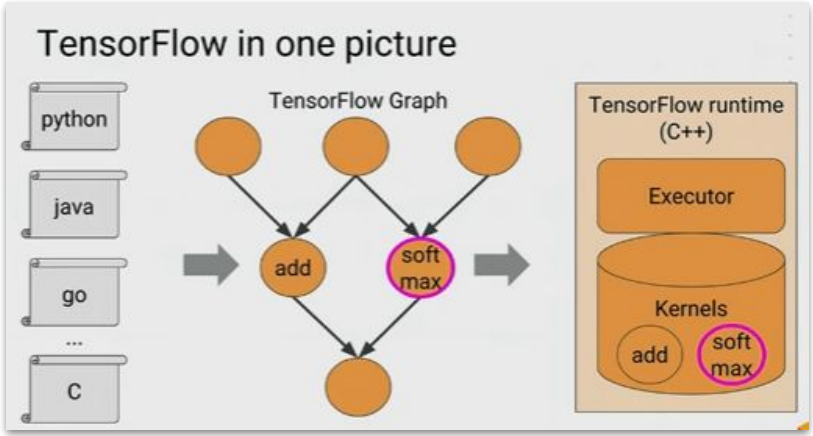

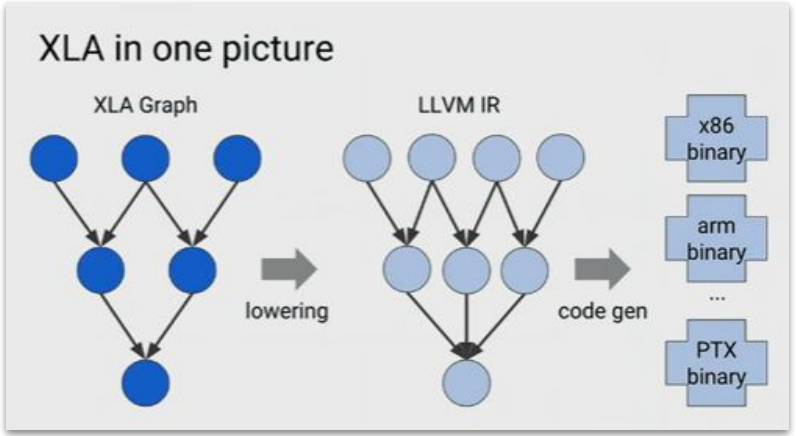

Optimized Implementation

Parallelism: Multi-thread, SIMD

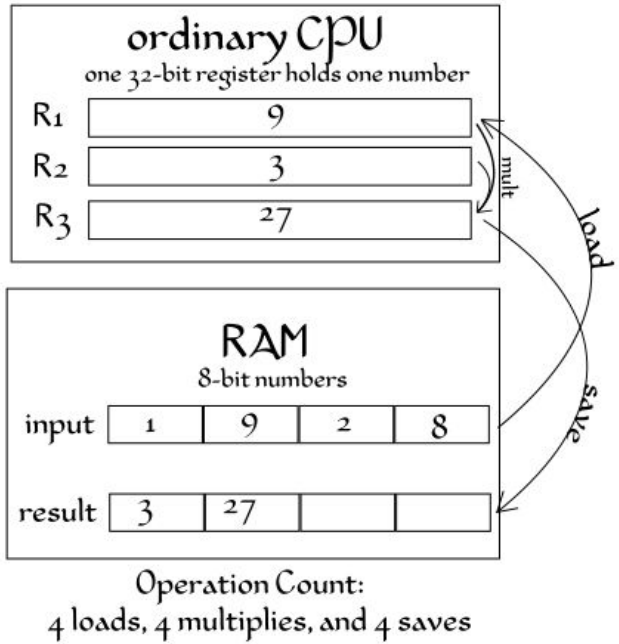

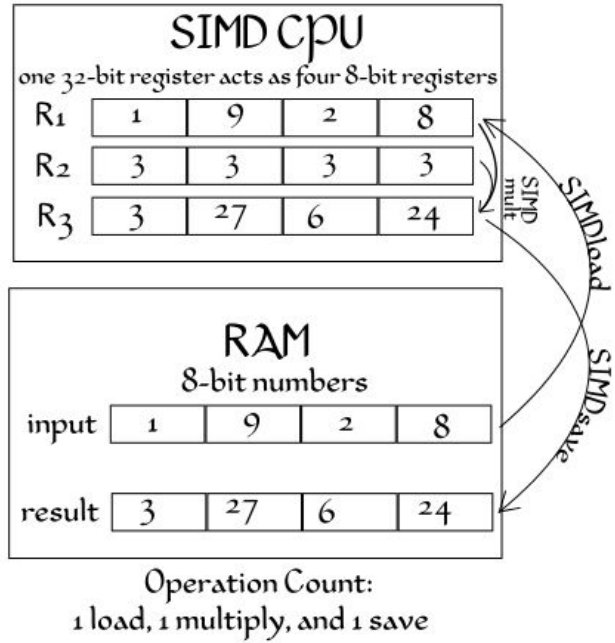

https://en.wikipedia.org/wiki/SIMD

Optimized Implementation

Data Locality

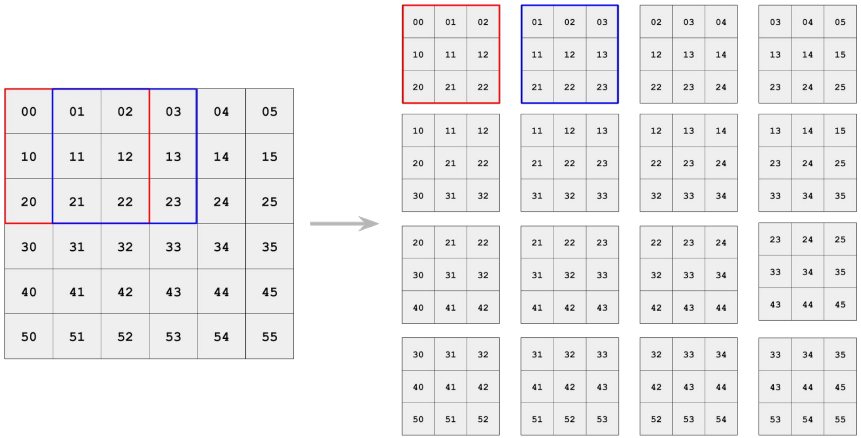

No Tiling

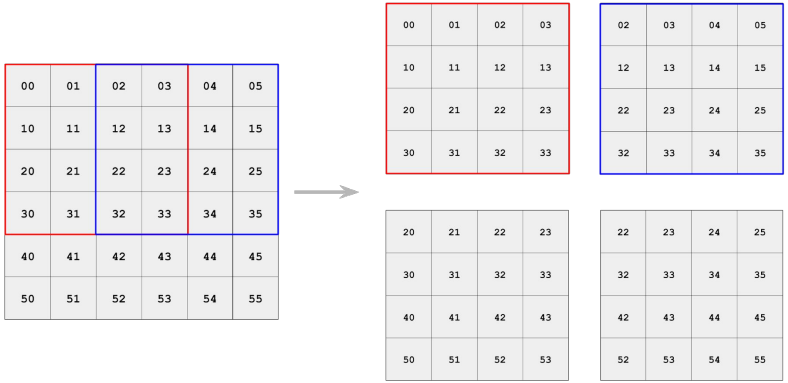

Tiling

https://tvm.ai/2017/08/22/Optimize-Deep-Learning-GPU-Operators-with-TVM-A-Depthwise-Convolution-Example.html

Optimized Implementation

Tradeoffs between parallelism, locality, and redundant computation

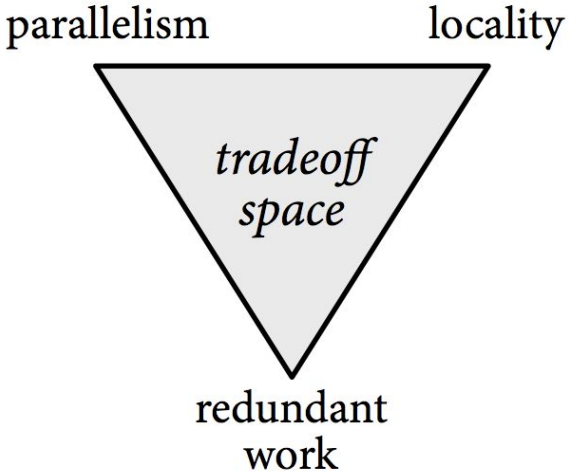

Ragan-Kelley J M. Decoupling algorithms from the organization of computation for high performance image processing[D]. Massachusetts Institute of Technology, 2014.

Optimized Implementation

Tools to optimize computation scheduling: Halide, TVM

(a) Clean C++ : 9.94 ms per megapixel

```
void blur(const Image &in, Image &blurred) {
 Image tmp(in.width(), in.height());

 for (int y = 0; y < in.height(); y++)
  for (int x = 0; x < in.width(); x++)
   tmp(x, y) = (in(x-1, y) + in(x, y) + in(x+1, y))/3;

 for (int y = 0; y < in.height(); y++)
  for (int x = 0; x < in.width(); x++)
   blurred(x, y) = (tmp(x, y-1) + tmp(x, y) + tmp(x, y+1))/3;
}
```

(b) Fast C++ (for x86) : 0.90 ms per megapixel

```
void fast_blur(const Image &in, Image &blurred) {
 __m128i one_third = _mm_set1_epi16(21846);
 #pragma omp parallel for
 for (int yTile = 0; yTile < in.height(); yTile += 32) {
  __m128i a, b, c, sum, avg;
  __m128i tmp[(256/8)*(32+2)];
  for (int xTile = 0; xTile < in.width(); xTile += 256) {
   __m128i *tmpPtr = tmp;
   for (int y = -1; y < 32+1; y++) {
    const uint16_t *inPtr = &(in(xTile, yTile+y));
    for (int x = 0; x < 256; x += 8) {
     a = _mm_loadu_si128((__m128i*)(inPtr-1));
     b = _mm_loadu_si128((__m128i*)(inPtr+1));
     c = _mm_load_si128((__m128i*)(inPtr));
     sum = _mm_add_epi16(_mm_add_epi16(a, b), c);
     avg = _mm_mulhi_epi16(sum, one_third);
     _mm_store_si128(tmpPtr++, avg);
     inPtr += 8;
   }}
   tmpPtr = tmp;
   for (int y = 0; y < 32; y++) {
    __m128i *outPtr = (__m128i *)(&(blurred(xTile, yTile+y)));
    for (int x = 0; x < 256; x += 8) {
     a = _mm_load_si128(tmpPtr+(2*256)/8);
     b = _mm_load_si128(tmpPtr+256/8);
     c = _mm_load_si128(tmpPtr++);
     sum = _mm_add_epi16(_mm_add_epi16(a, b), c);
     avg = _mm_mulhi_epi16(sum, one_third);
     _mm_store_si128(outPtr++, avg);
}}}}}
```

(c) Halide : 0.90 ms per megapixel

```
Func halide_blur(Func in) {
 Func tmp, blurred;
 Var x, y, xi, yi;

 // The algorithm
 tmp(x, y) = (in(x-1, y) + in(x, y) + in(x+1, y))/3;
 blurred(x, y) = (tmp(x, y-1) + tmp(x, y) + tmp(x, y+1))/3;

 // The schedule
 blurred.tile(x, y, xi, yi, 256, 32)
        .vectorize(xi, 8).parallel(y);
 tmp.chunk(x).vectorize(x, 8);

 return blurred;
}
```

Ragan-Kelley J, Adams A, Paris S, et al. Decoupling algorithms from schedules for easy optimization of image processing pipelines[J]. 2012.

Hardware Acceleration

GPU, FPGA, ASIC

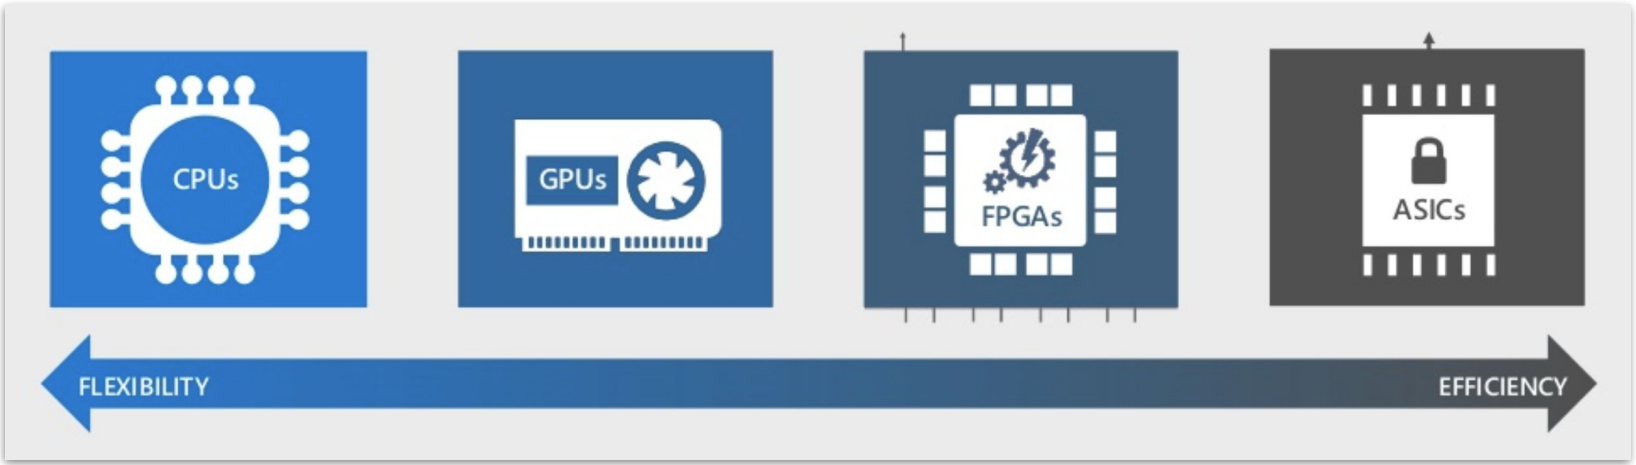

Mark Russinovich, “Inside Microsoft's FPGA-Based Configurable Cloud”, Microsoft Build 2017

Hardware Acceleration

Neural Processor Unit in Mobvoi A1

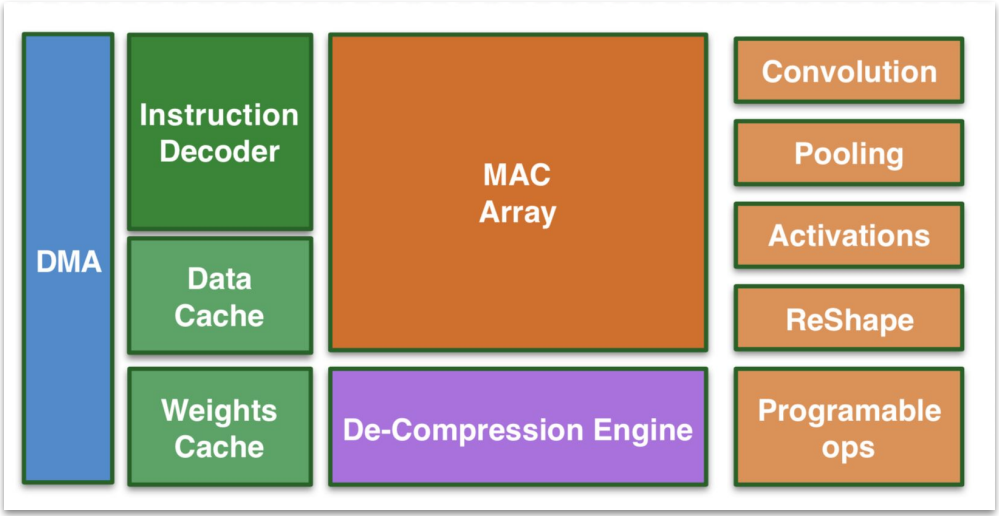

Neural Processor Unit (by NationalChip)

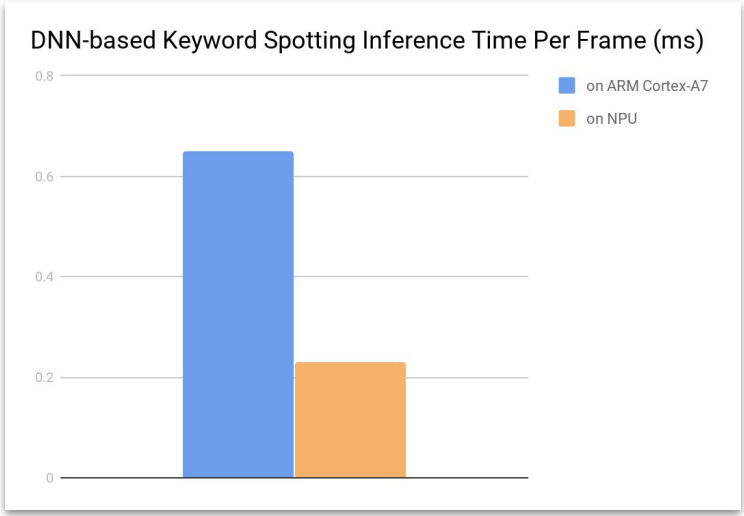

Inference on NPU achieve ~3x speed up compared with ARM Cortex-A7 Dual 1.2GHz CPU.

出门问问